

# Konstantinos (Kostas) Stavropoulos

Department of Computer Science, The University of Texas at Austin  
2317 Speedway, Austin, TX 78712, USA  
Email: [kstavrop@utexas.edu](mailto:kstavrop@utexas.edu) — Website: <https://www.kstavrop.com>

---

|                                                    |                                                                                                                                                                                                                             |                     |
|----------------------------------------------------|-----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|---------------------|
| RESEARCH INTERESTS                                 | Reliable Machine Learning; Computational Learning Theory.                                                                                                                                                                   |                     |
| EDUCATION                                          | <b>University of Texas at Austin</b><br>Ph.D. student, Computer Science<br><i>Advisor:</i> Adam Klivans                                                                                                                     | <i>2021–present</i> |
|                                                    | <b>National Technical University of Athens (NTUA)</b><br>Diploma in Electrical & Computer Engineering (5–year joint degree)<br><i>Thesis:</i> Learning rankings from incomplete samples<br><i>Advisor:</i> Dimitris Fotakis | <i>2015–20</i>      |
| SELECTED HONORS                                    | <b>Best Paper Award</b> , Foundations of Responsible Computing                                                                                                                                                              | <i>2026</i>         |
|                                                    | <b>Best Paper Award</b> , Reliable ML Workshop @ NeurIPS                                                                                                                                                                    | <i>2025</i>         |
|                                                    | <b>Dean’s Prestigious Fellowship Supplement</b> , UT Austin                                                                                                                                                                 | <i>2025</i>         |
|                                                    | <b>Apple Scholars in AI/ML PhD Fellowship</b>                                                                                                                                                                               | <i>2025</i>         |
|                                                    | <b>Best Paper Award</b> , Conference on Learning Theory                                                                                                                                                                     | <i>2024</i>         |
|                                                    | <b>Bodossaki, Leventis, and Gerondelis Fellowships</b>                                                                                                                                                                      | <i>2022–25</i>      |
| INDUSTRY EXPERIENCE                                | <b>State Scholarships Foundation Award</b> , Greece                                                                                                                                                                         | <i>2020</i>         |
|                                                    | <b>Apple Machine Learning Research</b><br>Research intern<br>with Parikshit Gopalan                                                                                                                                         | <i>Summer 2025</i>  |
| PUBLICATIONS<br>(author order $\alpha$ - $\beta$ ) | <b>21. Sandwiching Polynomials for Geometric Concepts with Low Intrinsic Dimension</b><br>Adam Klivans, Konstantinos Stavropoulos, Arsen Vasilyan<br><b>COLT 2026</b>                                                       |                     |
|                                                    | <b>20. Equivalence of Coarse and Fine-Grained Models for Learning with Distribution Shift</b><br>Adam Klivans, Shyamal Patel, Konstantinos Stavropoulos, Arsen Vasilyan<br><b>COLT 2026</b>                                 |                     |
|                                                    | <b>19. Testing Noise Assumptions of Learning Algorithms</b><br>Surbhi Goel, Adam Klivans, Konstantinos Stavropoulos, Arsen Vasilyan<br><b>Reliable ML Workshop @ NeurIPS 2025 ★ Best Paper ★</b><br><b>COLT 2026</b>        |                     |
|                                                    | <b>18. The Importance of Being Smoothly Calibrated</b><br>Parikshit Gopalan, Konstantinos Stavropoulos, Kunal Talwar, Pranay Tankala<br><b>FORC 2026 ★ Best Paper ★</b>                                                     |                     |
|                                                    | <b>17. Efficient Calibration for Decision Making</b><br>Parikshit Gopalan, Konstantinos Stavropoulos, Kunal Talwar, Pranay Tankala<br><b>STOC 2026</b>                                                                      |                     |
|                                                    | <b>16. A Fully Polynomial-Time Algorithm for Robustly Learning Halfspaces over the Hypercube</b><br>Gautam Chandrasekaran, Adam Klivans, Konstantinos Stavropoulos, Arsen Vasilyan<br><b>STOC 2026</b>                      |                     |
|                                                    |                                                                                                                                                                                                                             |                     |
|                                                    |                                                                                                                                                                                                                             |                     |

**15. Sparse Linear Regression is Easy on Random Supports**

Gautam Chandrasekaran, Raghu Meka, Konstantinos Stavropoulos  
**STOC 2026**

**14. The Power of Iterative Filtering for Supervised Learning with (Heavy) Contamination**

Adam Klivans, Konstantinos Stavropoulos, Kevin Tian, Arsen Vasilyan.  
**NeurIPS 2025 ★ Spotlight ★**

**13. Learning Constant-Depth Circuits in Malicious Noise Models**

Adam Klivans, Konstantinos Stavropoulos, Arsen Vasilyan  
**COLT 2025**

**12. Learning Neural Networks with Distribution Shift: Efficiently Certifiable Guarantees**

Gautam Chandrasekaran, Adam Klivans, Lin Lin Lee, Konstantinos Stavropoulos  
**ICLR 2025**

**11. Learning Noisy Halfspaces with a Margin: Massart is No Harder than Random**

Gautam Chandrasekaran, Vasilis Kontonis, Konstantinos Stavropoulos, Kevin Tian  
**NeurIPS 2024 ★ Spotlight ★**

**10. Tolerant Algorithms for Learning with Arbitrary Covariate Shift**

Surbhi Goel, Abhishek Shetty, Konstantinos Stavropoulos, Arsen Vasilyan  
**NeurIPS 2024 ★ Spotlight ★**

**9. Efficient Discrepancy Testing for Learning with Distribution Shift**

Gautam Chandrasekaran, Adam Klivans, Vasilis Kontonis, Konstantinos Stavropoulos, Arsen Vasilyan  
**NeurIPS 2024**

**8. Smoothed Analysis for Learning Concepts with Low Intrinsic Dimension**

Gautam Chandrasekaran, Adam Klivans, Vasilis Kontonis, Raghu Meka, Konstantinos Stavropoulos  
**COLT 2024 ★ Best Paper ★**

**7. Learning Intersections of Halfspaces with Distribution Shift: Improved Algorithms and SQ Lower Bounds**

Adam Klivans, Konstantinos Stavropoulos, Arsen Vasilyan  
**COLT 2024**

**6. Testable Learning with Distribution Shift**

Adam Klivans, Konstantinos Stavropoulos, Arsen Vasilyan  
**COLT 2024**

**5. An Efficient Tester-Learner for Halfspaces**

Aravind Gollakota, Adam Klivans, Konstantinos Stavropoulos, Arsen Vasilyan  
**ICLR 2024**

**4. Tester-Learners for Halfspaces: Universal Algorithms**

Aravind Gollakota, Adam Klivans, Konstantinos Stavropoulos, Arsen Vasilyan  
**NeurIPS 2023 ★ Oral ★**

**3. Agnostically Learning Single-Index Models using Omnipredictors**

Aravind Gollakota, Parikshit Gopalan, Adam Klivans, Konstantinos Stavropoulos  
**NeurIPS 2023**

**2. Learning and Covering Sums of Independent Random Variables with Unbounded Support**

Alkis Kalavasis, Konstantinos Stavropoulos, Manolis Zampetakis

NeurIPS 2022 ★ Oral ★

**1. Aggregating Incomplete and Noisy Rankings**

Dimitris Fotakis, Alkis Kalavasis, Konstantinos Stavropoulos

AISTATS 2021

PREPRINTS

**P1. Iterative Chow Filtering for Learning with Distribution Shift**

Gautam Chandrasekaran, George Gkrinias, Adam Klivans,  
Konstantinos Stavropoulos, Arsen Vasilyan

TALKS

**Efficient Learning Algorithms under (Heavy) Contamination**

MIT A&C Seminar

*Apr. 2026*

Georgia Tech ARC Colloquium

*Jan. 2026*

TTIC Young Researcher Seminar Series

*Nov. 2025*

MSR Postdoc Interview Talk

*Nov. 2025*

UPenn CS Theory Seminar

*Oct. 2025*

CMU CS Theory Lunch [\[video\]](#)

*Oct. 2025*

Apple MLR internal group meeting

*July 2025*

Stanford CS Theory Lunch

*July 2025*

**Efficiently Certifiable Guarantees for Learning with Distribution Shift**

Archimedes Center for Research in AI, Data Science and Algorithms

*Jan. 2025*

**Learning Constant-Depth Circuits in Malicious Noise Models**

UT Austin Pseudorandomness group meeting

*Nov. 2024*

**Learning Intersections of Halfspaces with Distribution Shift**

Conference on Learning Theory

*July 2024*

**Tester-Learners for Halfspaces: Universal Algorithms**

Oral Presentation, NeurIPS

*Dec. 2023*

**Learning and Covering Sums of Independent Random Variables with Unbounded Support**

Oral Presentation, NeurIPS

*Dec. 2022*

SERVICE &  
TEACHING

**Computational Foundations of Reliable Learning**

*June 2026*

Tutorial at COLT 2026, co-organized with Adam Klivans & Arsen Vasilyan

**Teaching Assistant, New Horizons Summer School in TCS**

*June 2023*

**Teaching Assistant, UT Austin**

*Course:* Principles of Machine Learning I: Honors

*Spring 2023*

*Instructor:* Adam Klivans

**Teaching Assistant, NTUA, Greece**

*Courses:* Algorithms and Complexity, Discrete Mathematics

*2020–2021*

*Instructor:* Dimitris Fotakis

**Reviewing/Subreviewing:** COLT (2025, 2026), FOCS (2025, 2026), STOC 2026, ITCS 2026, ICLR 2024, ICML 2024, NeurIPS (2023, 2026), RANDOM (2026)