

Konstantinos (Kostas) Stavropoulos

Department of Computer Science, The University of Texas at Austin
2317 Speedway, Austin, TX 78712, USA
Email: kstavrop@utexas.edu — Website: <https://www.kstavrop.com>

RESEARCH INTERESTS	Machine Learning; Robustness; Distribution Shift; Computational Learning Theory.	
EDUCATION	University of Texas at Austin Ph.D. student, Computer Science <i>Advisor:</i> Adam Klivans	<i>2021–present</i>
	National Technical University of Athens (NTUA) Diploma in Electrical & Computer Engineering (5–year joint degree) <i>GPA:</i> 9.76/10 (First in cohort) <i>Thesis:</i> Learning rankings from incomplete samples <i>Advisor:</i> Dimitris Fotakis	<i>2015–20</i>
AWARDS & FELLOWSHIPS	Best paper award, Reliable ML Workshop @ NeurIPS	<i>2025</i>
	Dean’s Prestigious Fellowship Supplement, UT Austin	<i>2025</i>
	Apple Scholars in AI/ML PhD fellowship	<i>2025</i>
	Best paper award, Conference on Learning Theory	<i>2024</i>
	External fellowships: Bodossaki, Leventis, Gerondelis	<i>2022–25</i>
	Scholarship award from Hellenic Professional Society of Texas	<i>2022</i>
	Award of Excellence from State Scholarships Foundation	<i>2020</i>
	Thomaideio Award from NTUA (highest GPA during a year)	<i>2019</i>
INDUSTRY EXPERIENCE	Award from Eurobank “The Great Moment for Education”	<i>2015</i>
	Apple Machine Learning Research Research intern with Parikshit Gopalan	<i>Summer 2025</i>
PUBLICATIONS	<i>(alphabetical author order)</i>	
	17. Efficient Calibration for Decision Making Parikshit Gopalan, Konstantinos Stavropoulos, Kunal Talwar, Pranay Tankala STOC 2026	
	16. A Fully Polynomial-Time Algorithm for Robustly Learning Halfspaces over the Hypercube Gautam Chandrasekaran, Adam Klivans, Konstantinos Stavropoulos, Arsen Vasilyan STOC 2026	
	15. Sparse Linear Regression is Easy on Random Supports Gautam Chandrasekaran, Raghu Meka, Konstantinos Stavropoulos STOC 2026	
	14. The Power of Iterative Filtering for Supervised Learning with (Heavy) Contamination Adam Klivans, Konstantinos Stavropoulos, Kevin Tian, Arsen Vasilyan. NeurIPS 2025 ★ Spotlight ★	

13. Learning Constant-Depth Circuits in Malicious Noise Models

Adam Klivans, Konstantinos Stavropoulos, Arsen Vasilyan

COLT 2025

12. Learning Neural Networks with Distribution Shift:

Efficiently Certifiable Guarantees

Gautam Chandrasekaran, Adam Klivans, Lin Lin Lee, Konstantinos Stavropoulos

ICLR 2025

11. Learning Noisy Halfspaces with a Margin:

Massart is No Harder than Random

Gautam Chandrasekaran, Vasilis Kontonis, Konstantinos Stavropoulos, Kevin Tian

NeurIPS 2024 ★ Spotlight ★

10. Tolerant Algorithms for Learning with Arbitrary Covariate Shift

Surbhi Goel, Abhishek Shetty, Konstantinos Stavropoulos, Arsen Vasilyan

NeurIPS 2024 ★ Spotlight ★

9. Efficient Discrepancy Testing for Learning with Distribution Shift

Gautam Chandrasekaran, Adam Klivans, Vasilis Kontonis,

Konstantinos Stavropoulos, Arsen Vasilyan

NeurIPS 2024

8. Smoothed Analysis for Learning Concepts with Low Intrinsic Dimension

Gautam Chandrasekaran, Adam Klivans, Vasilis Kontonis,

Raghu Meka, Konstantinos Stavropoulos

COLT 2024 ★ Best Paper ★

7. Learning Intersections of Halfspaces with Distribution Shift:

Improved Algorithms and SQ Lower Bounds

Adam Klivans, Konstantinos Stavropoulos, Arsen Vasilyan

COLT 2024

6. Testable Learning with Distribution Shift

Adam Klivans, Konstantinos Stavropoulos, Arsen Vasilyan

COLT 2024

5. An Efficient Tester-Learner for Halfspaces

Aravind Gollakota, Adam Klivans, Konstantinos Stavropoulos, Arsen Vasilyan

ICLR 2024

4. Tester-Learners for Halfspaces: Universal Algorithms

Aravind Gollakota, Adam Klivans, Konstantinos Stavropoulos, Arsen Vasilyan

NeurIPS 2023 ★ Oral ★

3. Agnostically Learning Single-Index Models using Omnipredictors

Aravind Gollakota, Parikshit Gopalan, Adam Klivans, Konstantinos Stavropoulos

NeurIPS 2023

**2. Learning and Covering Sums of Independent Random Variables
with Unbounded Support**

Alkis Kalavasis, Konstantinos Stavropoulos, Manolis Zampetakis

NeurIPS 2022 ★ Oral ★

	1. Aggregating Incomplete and Noisy Rankings Dimitris Fotakis, Alkis Kalavasis, Konstantinos Stavropoulos AISTATS 2021	
PREPRINTS	P1. Testing Noise Assumptions of Learning Algorithms Surbhi Goel, Adam Klivans, Konstantinos Stavropoulos, Arsen Vasilyan Reliable ML Workshop @ NeurIPS 2025 ★ Best Paper ★ [arXiv:2501.09189]	
TALKS	Efficient Learning Algorithms under (Heavy) Contamination Georgia Tech ARC Colloquium <i>Jan. 2026</i> TTIC Young Researcher Seminar Series <i>Nov. 2025</i> MSR Postdoc Interview Talk <i>Nov. 2025</i> UPenn CS Theory Seminar <i>Oct. 2025</i> CMU CS Theory Lunch [video] <i>Oct. 2025</i> Apple MLR internal group meeting <i>July 2025</i> Stanford CS Theory Lunch <i>July 2025</i> Efficiently Certifiable Guarantees for Learning with Distribution Shift Archimedes Center for Research in AI, Data Science and Algorithms <i>Jan. 2025</i> Learning Constant-Depth Circuits in Malicious Noise Models UT Austin Pseudorandomness group meeting <i>Nov. 2024</i> Learning Intersections of Halfspaces with Distribution Shift Conference on Learning Theory <i>July 2024</i> Tester-Learners for Halfspaces: Universal Algorithms Oral Presentation, NeurIPS <i>Dec. 2023</i> Learning and Covering Sums of Independent Random Variables with Unbounded Support Oral Presentation, NeurIPS <i>Dec. 2022</i>	
SERVICE & TEACHING	Teaching Assistant, New Horizons Summer School in TCS <i>June 2023</i> Teaching Assistant, UT Austin <i>Course:</i> Principles of Machine Learning I: Honors <i>Spring 2023</i> <i>Instructor:</i> Adam Klivans Teaching Assistant, NTUA, Greece <i>Courses:</i> Algorithms and Complexity, Discrete Mathematics <i>2020–2021</i> <i>Instructor:</i> Dimitris Fotakis Reviewing: ITCS 2026, FOCS 2025, COLT 2025, ICLR 2024, ICML 2024, NeurIPS 2023	
LANGUAGES AND SKILLS	English (fluent), French (basic), Greek (native) Python, $\text{\LaTeX}$, C/C++	