Konstantinos (Kostas) Stavropoulos

Department of Computer Science, The University of Texas at Austin
2317 Speedway, Austin, TX 78712, USA
Email: kstavrop@utexas.edu — Website: https://www.kstavrop.com

---

Research Interests

Machine Learning; Robustness; Distribution Shift; Computational Learning Theory.

Education

University of Texas at Austin *2021–present*
Ph.D. student, Computer Science
Advisor: Adam Klivans

National Technical University of Athens (NTUA) *2015–20*
Diploma in Electrical & Computer Engineering (5–year joint degree)
GPA: 9.76/10 (First in cohort)
Thesis: Learning rankings from incomplete samples
Advisor: Dimitris Fotakis

Awards & Fellowships

Best paper award, Reliable ML Workshop @ NeurIPS *2025*
Apple Scholars in AI/ML PhD fellowship *2025*
Best paper award, Conference on Learning Theory *2024*
External fellowships: Bodossaki, Leventis, Gerondelis *2022–25*
Scholarship award from Hellenic Professional Society of Texas *2022*
Award of Excellence from State Scholarships Foundation *2020*
Thomaideio Award from NTUA (highest GPA during a year) *2019*
Award from Eurobank "The Great Moment for Education" *2015*

Industry Experience

Apple Machine Learning Research *Summer 2025*
Research intern
with Parikshit Gopalan

Publications

(alphabetical author order)

14. The Power of Iterative Filtering for Supervised Learning with (Heavy) Contamination
Adam Klivans, Konstantinos Stavropoulos, Kevin Tian, Arsen Vasilyan.
NeurIPS 2025 ★ Spotlight ★

13. Learning Constant-Depth Circuits in Malicious Noise Models
Adam Klivans, Konstantinos Stavropoulos, Arsen Vasilyan
COLT 2025

12. Learning Neural Networks with Distribution Shift: Efficiently Certifiable Guarantees
Gautam Chandrasekaran, Adam Klivans, Lin Lin Lee, Konstantinos Stavropoulos
ICLR 2025

11. Learning Noisy Halfspaces with a Margin: Massart is No Harder than Random
Gautam Chandrasekaran, Vasilis Kontonis, Konstantinos Stavropoulos, Kevin Tian
NeurIPS 2024 ★ Spotlight ★

10. Tolerant Algorithms for Learning with Arbitrary Covariate Shift
Surbhi Goel, Abhishek Shetty, Konstantinos Stavropoulos, Arsen Vasilyan
NeurIPS 2024 ★ Spotlight ★

9. Efficient Discrepancy Testing for Learning with Distribution Shift
Gautam Chandrasekaran, Adam Klivans, Vasilis Kontonis,
Konstantinos Stavropoulos, Arsen Vasilyan
NeurIPS 2024

8. Smoothed Analysis for Learning Concepts with Low Intrinsic Dimension
Gautam Chandrasekaran, Adam Klivans, Vasilis Kontonis,
Raghu Meka, Konstantinos Stavropoulos
COLT 2024 ★ Best Paper ★

7. Learning Intersections of Halfspaces with Distribution Shift: Improved Algorithms and SQ Lower Bounds
Adam Klivans, Konstantinos Stavropoulos, Arsen Vasilyan
COLT 2024

6. Testable Learning with Distribution Shift
Adam Klivans, Konstantinos Stavropoulos, Arsen Vasilyan
COLT 2024

5. An Efficient Tester-Learner for Halfspaces
Aravind Gollakota, Adam Klivans, Konstantinos Stavropoulos, Arsen Vasilyan
ICLR 2024

4. Tester-Learners for Halfspaces: Universal Algorithms
Aravind Gollakota, Adam Klivans, Konstantinos Stavropoulos, Arsen Vasilyan
NeurIPS 2023 ★ Oral ★

3. Agnostically Learning Single-Index Models using Omnipredictors
Aravind Gollakota, Parikshit Gopalan, Adam Klivans, Konstantinos Stavropoulos
NeurIPS 2023

2. Learning and Covering Sums of Independent Random Variables with Unbounded Support
Alkis Kalavasis, Konstantinos Stavropoulos, Manolis Zampetakis
NeurIPS 2022 ★ Oral ★

1. Aggregating Incomplete and Noisy Rankings
Dimitris Fotakis, Alkis Kalavasis, Konstantinos Stavropoulos
AISTATS 2021

Preprints

P4. Efficient Calibration for Decision Making
Parikshit Gopalan, Konstantinos Stavropoulos, Kunal Talwar, Pranay Tankala
[arXiv:2511.13699]

P3. A Fully Polynomial-Time Algorithm for Robustly Learning Halfspaces over the Hypercube
Gautam Chandrasekaran, Adam Klivans, Konstantinos Stavropoulos, Arsen Vasilyan
[arXiv:2511.07244]

P2. Sparse Linear Regression is Easy on Random Supports
Gautam Chandrasekaran, Raghu Meka, Konstantinos Stavropoulos
[arXiv:2511.06211]

P1. Testing Noise Assumptions of Learning Algorithms
Surbhi Goel, Adam Klivans, Konstantinos Stavropoulos, Arsen Vasilyan
Reliable ML Workshop @ NeurIPS 2025 ★ **Best Paper** ★
[arXiv:2501.09189]

Talks

Efficient Learning Algorithms under (Heavy) Contamination
TTIC Young Researcher Seminar Series, MSR Postdoc Interview Talk — *Nov. 2025*
UPenn CS Theory Seminar, CMU CS Theory Lunch [video] — *Oct. 2025*
Apple MLR internal group meeting, Stanford CS Theory Lunch — *July 2025*

Efficiently Certifiable Guarantees for Learning with Distribution Shift
Archimedes Center for Research in AI, Data Science and Algorithms — *Jan. 2025*

Learning Constant-Depth Circuits in Malicious Noise Models
UT Austin Pseudorandomness group meeting — *Nov. 2024*

Learning Intersections of Halfspaces with Distribution Shift
Conference on Learning Theory — *July 2024*

Tester-Learners for Halfspaces: Universal Algorithms
Oral Presentation, NeurIPS — *Dec. 2023*

Learning and Covering Sums of Independent Random Variables with Unbounded Support
Oral Presentation, NeurIPS — *Dec. 2022*

Service & Teaching

Teaching Assistant, New Horizons Summer School in TCS — *June 2023*

Teaching Assistant, UT Austin
Course: Principles of Machine Learning I: Honors — *Spring 2023*
Instructor: Adam Klivans

Teaching Assistant, NTUA, Greece
Courses: Algorithms and Complexity, Discrete Mathematics — *2020–2021*
Instructor: Dimitris Fotakis

Reviewing: ITCS 2026, FOCS 2025, COLT 2025, ICLR 2024, ICML 2024, NeurIPS 2023

Languages and Skills

English (fluent), French (basic), Greek (native)
Python, LaTeX, C/C++